



CLOSE THE VISIBILITY GAP

What Delayed Integration Costs Your Healthcare Portfolio

Executive Summary

The success of PE-backed healthcare rollups depends on how fast a platform can unify operations after each acquisition. Most cannot do it fast enough. Each acquired practice brings its own EHR, billing system, and years of accumulated data conventions. The conventional approach gets more expensive with every new acquisition, not less:

- **Visibility gap cost.** For a mid-to-large platform, the combined monthly cost of incomplete integration runs \$200,000 to \$500,000 in revenue leakage, duplicative labor, and legacy system carrying costs.
- **Compounding maintenance.** Every new acquisition makes existing integrations more expensive. At fifty, the team is consumed by maintenance. New acquisitions stall.
- **Hold period drag.** Twelve to eighteen months lost to integration means twenty to thirty percent of a five-year hold producing no synergies. That is capital deployed but not yet compounding.
- **Synergy failure.** Nearly seventy percent of mergers fail to achieve expected revenue synergies, and the integration approach most platforms use is a primary driver.

MTN's Data Foundry compresses integration timelines from months to weeks with a team of three to five instead of fifteen to twenty-five, and keeps maintenance flat as the portfolio grows. For a ten-clinic portfolio, that acceleration unlocks over **\$2 million in year-one EBITDA**, translating to approximately **\$24 million in enterprise value** at a 12x exit multiple.

Intended Audience (Executive Edition)

Target Audience: CEOs, CFOs, COOs, and Operating Partners at PE-backed healthcare platforms

Business Focus: EBITDA acceleration, exit readiness, and hold-period efficiency through faster data integration.

Key Economics: \$208K additional EBITDA per clinic per year from integration acceleration; \$24M enterprise value at 12x across 10 acquisitions.

Comparison: 15–25 people over 8–10 weeks compressed to 3–5 people in 2–3 weeks with flat maintenance at scale.

The Cost of Delayed Integration: Why Every Acquisition Makes It Worse

A PE-backed dermatology platform closes on its fifth acquisition in eighteen months. Each practice arrived with its own EHR, billing system, and data conventions. The integration consultants are four months into the third site and have not started the fourth. The CFO cannot produce a consolidated denial report for the board. Meanwhile, the sixth deal is already in letter of intent.

This is the default experience of healthcare consolidation in the United States today. And the gap between the pace of acquisition and the pace of data integration is the single greatest underappreciated risk in PE-backed healthcare rollups.

THE CONSOLIDATION WAVE

Add-on acquisitions now account for approximately seventy-three percent of all PE buyout deals (Choi, Jinny and Walters, Kyle and Tang, Kenny and Corridore, Jim and Wright, Brian, 2025; CohnReznick, 2024). The target specialties, dermatology, ophthalmology, gastroenterology, dental, orthopedics, and behavioral health, share the profile that makes rollups attractive: fragmented markets, favorable payer mixes, and high-margin procedural revenue. They also share the profile that makes data integration hardest, because no two practices encode their operations the same way.

The economic logic depends entirely on integration speed. Top-performing PE firms target operational unification within six months of close, because every month of delayed integration is a month of unrealized synergies: duplicate billing teams, fragmented reporting, and missed opportunities to optimize across the portfolio.

WHAT THE VISIBILITY GAP COSTS

Every month of incomplete integration extends a visibility gap during which consolidated analytics and revenue cycle improvements simply do not exist. Five clinics, five different ways of recording why a claim was denied. One uses industry-standard codes. Another uses free text. A third tracks only paid-or-not. Until someone translates all five into a common language, the CFO cannot tell the board what the top denial reasons are across the platform.

For a mid-to-large platform, the combined monthly cost of this visibility gap runs **\$200,000 to \$500,000**, spanning three categories: revenue leakage from unresolved denials, duplicative labor across parallel workflows, and the carrying cost of legacy systems that cannot yet be decommissioned. Over forty percent of healthcare organizations report losing ten percent or more of annual revenue to leakage (Sage Growth Partners and Fibroblast, 2024), and up to sixty-five percent of denied claims are never resubmitted (Healthcare Financial Management Association, 2024).

How many months after your last close before the CFO could produce a consolidated denial report for the board? If the answer is “we are still working on it,” you are operating the portfolio on faith rather than data.

WHY IT GETS WORSE WITH EVERY ACQUISITION

The eighth acquisition should be easier than the first, not harder. But with conventional integration, every new system you add makes every existing integration more expensive to maintain. The reason is structural: changes in one system’s data formats, vendor updates, coding standards, or local conventions can break integrations with every other system already connected.

At fifty active integrations, the team spends all its capacity maintaining what already exists. New acquisitions stall. The rollup thesis breaks.

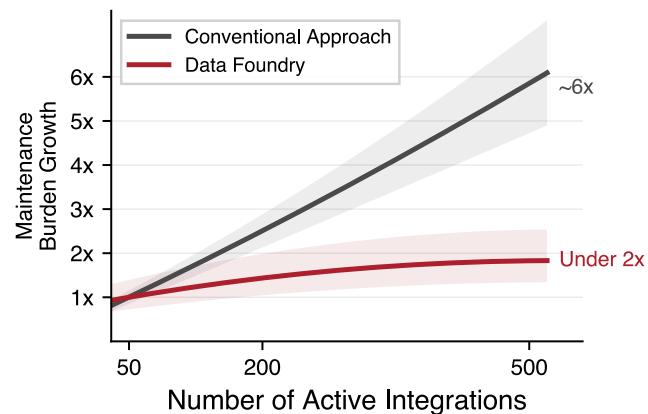


Figure 1. As portfolio size grows, conventional integration maintenance compounds. Each new system increases the burden on every existing integration. A shared-layer approach keeps maintenance growth nearly flat. Shaded regions indicate ranges based on portfolio complexity.

A typical enterprise healthcare data integration project runs approximately \$2 million, with a team of fifteen to twenty-five people working eight to ten weeks per batch (Ganti, Aditya, 2025). The conventional approach resets to near-zero with each acquisition. Nothing compounds. Nothing scales.

THE HOLD PERIOD PROBLEM

A five-year hold that loses twelve to eighteen months to integration is really a three-and-a-half to four-year hold. That is twenty to thirty percent of the period your investors are paying for, spent waiting for systems to talk to each other instead of producing returns. Manual data processes alone extend deal timelines by an estimated thirty percent (Ganti, Aditya, 2025).

The project risk is equally real. Large-scale IT projects run an average of forty-five percent over budget while delivering fifty-six percent less value than predicted. One in six becomes a black swan with cost overruns exceeding two hundred percent (McKinsey & Company, 2012). A CEO or CFO who has lived through a failed ERP implementation understands this risk. The question is not whether integration is necessary. It is whether you can afford to do it the conventional way.

Compressing Integration from Months to Weeks: A Platform Built for Serial Acquirers

WHAT IT DOES

MTN's Data Foundry platform takes a fundamentally different approach to healthcare data integration. Rather than building connections between every system and every other system, the platform maps each source to a single shared layer. The result: adding the fiftieth source does not require touching sources one through forty-nine.

The system automatically proposes data mappings for human review and approval. An analyst, not a senior engineer, can review and approve a typical integration in minutes rather than days. When source systems change, the platform detects the deviation automatically and queues it for review. This is what keeps the maintenance burden flat regardless of portfolio size.

In benchmark testing across the structural diversity found in healthcare rollups, the system achieves over

ninety-nine percent mapping accuracy. Anything the system is not confident about is flagged for human review. Nothing goes live without explicit sign-off.

THE COMPARISON

Your current approach requires fifteen to twenty-five people over eight to ten weeks per integration batch. Data Foundry does it with three to five people in two to three weeks.

At fifty active integrations, the conventional approach requires two to three dedicated engineers just to keep existing integrations running. Data Foundry can substantially reduce that burden.

A typical enterprise integration project at this scale runs into the millions, with large teams working months per batch. The platform compresses both cost and timeline by eliminating the structural redundancy that makes conventional integration expensive in the first place.

EBITDA IMPACT

Consider a newly acquired practice generating \$10 million in annual revenue at a fifteen percent EBITDA margin. The PE thesis projects that centralized billing, improved denial management, and operational economies will raise that margin to twenty percent, producing an additional \$500,000 in annual EBITDA. Traditional integration delays those improvements by six months or more. By compressing integration to one to two months, roughly four to five additional months of enhanced margin are realized in year one: approximately **\$208,000 in additional EBITDA per clinic**.

Across a portfolio of ten acquisitions, the acceleration effect alone exceeds **\$2 million in year-one EBITDA**. At a representative exit multiple of 12x, that translates to approximately **\$24 million in enterprise value**. A return that dwarfs the cost of the integration platform itself.

EXIT READINESS

At exit, buyers pay more for genuinely integrated platforms than for a collection of logos under a holding company. Practitioner commentary frames this directly: "real integration in buy-and-builds" avoids the discount that comes with incomplete unification, while genuinely integrated platforms earn multiple expansion (Cohn-Reznick, 2024). Even a 0.5x to 1.0x multiple differential

on exit EBITDA can be worth tens of millions.

Firms that invest in modern data architecture outperform peers by fifteen to twenty percent in portfolio company value creation (Ganti, Aditya, 2025). Integration quality is not a back-office concern. It is exit defense.

Next Steps

The conventional integration model does not scale. The question is what you do about it.

INTEGRATION COMPLEXITY ASSESSMENT

We offer an integration complexity assessment: a two-week analysis that maps your current data landscape, identifies the specific visibility gaps across your portfolio, and estimates the cost of delayed integration at your current acquisition pace. To schedule an assessment, contact warren@themtn.ai.

Whether you are evaluating a platform acquisition, midway through a multi-site integration, or planning for the next phase of growth, the core question is the same: how many active data sources will you have in eighteen months, and does your current integration approach scale to that number?

MTN is a research and technology company with deep roots in clinical operations and machine intelligence. We bring to these conversations the perspective of researchers and advisors with clinical, technical, and health policy backgrounds.

TECHNICAL LEADERSHIP

Our team's work has been published in Nature journals, PNAS, JMIR, Chest, PLoS Computational Biology, The Royal Society, and other leading venues.

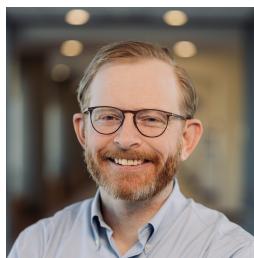


**Warren Pettine, MD,
Co-Founder and CEO.**

Assistant Professor at the University of Utah where he leads the Medical Machine Intelligence (M²Int) Lab. Trained in machine learning research at Harvard, Stanford, NYU, and Yale. Prior health policy experience in the U.S. Congress and service on the University of Utah Institutional Review Board ground MTN's approach in policy and regulatory expertise.



Matthias Christenson, PhD, AI Architect. Investigator with the M²Int Lab. PhD and postdoctoral research at Columbia University in computational ML, with prior industry experience as a Deep Learning Research Engineer at DeepLife training foundational models on genomic and biometric data. Leads MTN's technical architecture design and data model development.



Brian Locke, MD, MSCI, Clinical AI Lead. Investigator with the M²Int Lab. Active ICU physician and Assistant Professor at Intermountain Healthcare, bringing firsthand understanding of clinical workflows across academic medical centers and integrated delivery networks. Provides the methodological rigor for the clinical and operational implications of MTN's technology.



Samuel Wecker, Lead Systems Engineer. Over twelve years building and scaling production software, including as a founding engineer at a startup that grew to a billion-dollar platform. Specializes in unifying disparate systems and data sources at scale. Leads Data Foundry's core platform development.

Contact:

Warren Pettine, MD, CEO of MTN

warren@themtn.ai

<https://www.themtn.ai>