



TURN INTEGRATION LABOR INTO COMPOUNDING IP

Build the Accelerator Once. Compound It Across Every PE Bolt-On.

Executive Summary

The senior consultants on your healthcare M&A team are exceptionally good at integrating disparate clinical and billing systems across PE-backed bolt-ons. The work is profitable, the methodology is real, and almost none of it carries to the next acquisition. Intellectual Property that should be accruing to your firm is being rebuilt by hand and walked out the door with each engagement.

The MTN Data Foundry augments your team's work with AI-assisted mapping that builds a semantic data model across sources. The Foundry operates on schema metadata: table and field names, data types, API documentation, internal data dictionaries, and your team's accumulated knowledge. Patient and financial data values stay in your client's environment; the platform never sees them.

- **Your talent focuses on the high-value work.** Schema discovery and mapping that conventionally consumes 6 to 8 weeks compresses to 1 to 2 weeks during the pilot.
- **Maintain your data model, even as the software changes.** The Foundry automatically detects drift and suggests new mappings.
- **Every engagement deposits IP into your firm's library.** The Foundry captures three reusable assets on each deal: the annotation conventions your seniors apply, the firm's own documented mapping knowledge, and a growing library of canonical concepts (Patient.MRN, Claim.DenialReason, Provider.NPI) that every future source maps against. Adding the tenth EHR does not require touching the first nine. By engagement five, the marginal cost of integration is reviewing a queue, not building from scratch.

Intended Audience (Design Partner Briefing)

Target Audience: Healthcare M&A practice leaders evaluating a tool to deploy on the next PE bolt-on integration engagement.

Business Focus: A 90-day Design Partner pilot for the MTN Data Foundry. Single engagement, no exclusivity, mutually opt-in continuation.

Key Question: Can a productized integration platform cut our next engagement's data work and produce reusable templates for the bolt-on after that?

Comparison: Today's manual schema mapping and brittle pipelines vs. an AI-assisted data fabric with healthcare schema templates.




Figure 1. Days to integrated, reportable data per bolt-on under the conventional approach (flat at approximately three months) and under MTN Foundry. Two effects: engagement one is compressed by healthcare templates and automated mapping (no compounding required); engagements two and onward compress further as cross-engagement reuse caches mapping decisions, declining toward a 14-day asymptote by the fifth bolt-on. Forward-looking model with conservative calibration.

Your Team's Talent is Under-Leveraged

Your senior consultants know Epic Clarity, athenahealth, and eClinicalWorks at a level few in healthcare can match. They read clinical workflows and payer-segment ambiguity that no generalist consultancy gets right. That expertise is what your firm's reputation rests on.

But healthcare M&A integration is a slog. Three EHRs, two PM systems, a stale data warehouse, a PE sponsor's MD asking for payer mix across eight practices by Friday. The pipelines you build break silently when an EHR upgrade renames a column, and none of the work carries to the next bolt-on.

The PE sponsor's MD wants payer mix, days in AR, denial reasons, and provider productivity, comparable across the platform, refreshed before the next LP letter. Nothing about that ask is new. What is new is that it has to land eight times after eight closes, then nine, then ten, on a clock your firm does not control.

The Foundry takes the mechanical mapping layer off your senior consultants' desks: schema discovery, field-level annotation, and the canonical concept map that lets one source plug into a shared semantic layer instead of into every other source pairwise. The judg-

ment work stays human (architecture, exception handling, clinical-fidelity calls), and value-level normalization stays where it belongs, in your firm's methodology and your client's data environment. The methodology your team applies is encoded once, audited, and carries into the next engagement.

A SOLUTION THAT CONVERGES

For the practice leader. Concurrency lift on every senior consultant, fixed-fee margin lift on every engagement that can shift off T&M, and a productized accelerator that survives consultant turnover.

For your senior integration architect. The mechanical mapping work disappears; architecture, exception handling, and clinical-fidelity judgment stay human. Confidence scoring, audit logs, and the canonical schema model are visible and controllable.

For your engagement partner with the named PE sponsor. A faster commitment your competitors cannot credibly match. Your team owns the client relationship and the deliverable; MTN engineering deploys in support, not in front, ensuring the platform performs against the engagement's specifics.

How the Foundry Works

The MTN Data Foundry is an AI-assisted data integration and normalization platform. Given a new source, it reasons over schema, software documentation, and your team's accumulated knowledge to discover the common data model and the mapping decisions that get there, with confidence scoring and human review on every key step.

The Foundry's integration technology operates on meta-data: schema names, data types, API documentation, internal data dictionaries, and your team's accumulated knowledge. Patient and financial data values stay in your client's environment; the platform never sees them. This eliminates, rather than reduces, the PHI handling overhead that blocks most integration projects before they start. On top of that boundary, three AI-assisted capabilities replace the manual labor that consumes most of a conventional integration:

- **Annotations describing the function and intent of every table and field.** Your team spends hours painstakingly hunting down the meaning of columns in important but isolated tables. The MTN Data Foundry reasons over database structure, software documentation, organizational documents and your team's accumulated knowledge to annotate the database. Annotations are suggested by the system, flagged with confidence levels for human review. These can be used both by individuals, and super-power any software that uses the data.
- **Flexibly join data across sources.** Joining data across multiple sources involves manually wrangling tables from one form to another. Rather than building pairwise mappings between every source and every other source (the $N \times M$ explosion that consumes 65 to 75% of integration labor), the Foundry maps every source to a shared concept layer. Adding the tenth EHR does not require touching the first nine; maintenance stays manageable as the portfolio grows.
- **Self-healing maintenance.** When an upstream EHR adds or renames a field, the Foundry detects the deviation, proposes the updated mapping, and queues it for review rather than letting the integration break silently weeks later.

The platform is bounded to the schema-and-concept layer by design. Value-level normalization (applying ICD-9 to ICD-10 crosswalks, reconciling enum vocabularies, harmonizing payer-name variants) is a separately-deployable module that runs inside your client's environment when scope requires it. This white paper focuses on the schema-and-concept layer because that is where senior-consultant time is consumed and where productized methodology compounds across engagements.

What the Accelerator Is

An accelerator is an asset with pre-built tooling, methodology stacks, and AI-enabled platforms that layer into M&A integration projects. Every engagement you enter draws on the accumulated knowledge of engagements past. With the MTN Data Foundry, that asset has several constituent parts:

1. **Annotation conventions, your firm's house style, encoded.** Every engagement, senior consultants confirm or edit field-level annotations: usage notes, display names, definitions, and the reasoning behind each mapping. These are structured, versioned artifacts with confidence scores attached. The next engagement starts pre-loaded with the firm's accumulated conventions, and the suggestion engine is calibrated against the firm's own gold-standard reviews.
2. **Internal documented knowledge, made operational.** The Foundry reasons over your firm's own integration playbooks, methodology decks, prior engagement deliverables, and internal data dictionaries, using them as context for mapping reasoning so first-pass annotations reflect your documented approach rather than a generic prior. Document libraries that previously lived dormant in SharePoint become active mapping intelligence on every engagement.
3. **Canonical concepts derived from past engagements.** A growing library of stable business objects (`Patient.MRN`, `Encounter.EncounterType`, `Claim.DenialReason`, `Provider.NPI`, `Practice.RVU`) that every future source maps against. By engagement five, the library has crystallized around your firm's vertical (MSO consolidation, dental DSO, behavioral health, RCM rollups), and the marginal cost of

the next bolt-on falls toward the cost of reviewing a queue.

4. **Vendor-schema drift signatures and reviewer feedback loops.** When an upstream EHR renames a field in a quarterly release, the Foundry's correction is captured and applies firm-wide. Every approve, reject, or edit decision a senior consultant makes becomes labeled training data calibrating the system for your firm's standards.

The MTN Data Foundry accelerates work during the pilot period, and value compounds with every subsequent engagement. Pilot benchmarks suggest a senior consultant who today delivers 4 to 6 PE bolt-on integrations per year can deliver materially more on a productized platform, at higher margin per deal. We will commit to specific targets jointly with our design partner based on engagement scope.

WHAT ENGAGEMENT ONE LOOKS LIKE

The conventional first integration on a typical PE bolt-on takes roughly three months. The Foundry compresses that work substantially on engagement one, before any cross-engagement compounding has kicked in. Engagement one delivers a canonical concept map across the in-scope source systems and a queryable semantic layer your team can build the PE-sponsor deliverable on top of. MTN's engineering talent deploys in support of your team, ensuring the Data Foundry delivers.

What Success Looks Like

A successful pilot produces:

- Detailed annotations for every table and field, facilitating both general business intelligence and the use of advanced AI tooling.
- A clear mapping of data sources to either a target schema, or canonical set of concepts.
- A unified payer-mix and KPI dashboard delivered to the PE sponsor, built on the canonical concept map the Foundry produces, with your firm's normalization and clinical-fidelity rules applied in your delivery environment.
- An audit log of every transform applied.

After the pilot, the relationship is mutually opt-in. Whether you continue or not, you keep your annotation

conventions, methodology playbooks, vertical templates, and engagement deliverables, exported as portable JSON manifests. MTN keeps the platform and the learnings about how it performs across deployments. License terms, support model, and account scoping for an ongoing relationship are documented in the design partner term sheet, available on request.

What We Are Asking For

We are seeking our first healthcare-M&A integration design partner. The terms below reflect that: discounted, scope-bounded, and structured to build a case study together.

Three concrete asks over the next 90 days.

1. A 30-minute scoping conversation with the M&A practice leader for healthcare and the engagement lead for an upcoming PE bolt-on, to confirm whether the Foundry's current EHR and PM coverage matches the engagement.
2. Identification of a candidate engagement in your firm's pipeline (MSO consolidation, dental services, behavioral health, ambulatory surgery, RCM rollups, payer integration) that could host the 90-day pilot.
3. A 90-day Design Partner agreement covering platform license discount, engineering support and other key details. Draft term sheet ready for review.

No exclusive in either direction.

For the platform's technical architecture and compliance posture, request the *MTN Data Foundry Architecture and Compliance Brief*.

About MTN

MTN is a team with deep expertise in healthcare data infrastructure, machine learning, and clinical informatics. The founding team includes a physician-researcher and experienced engineers who have built semantic mapping systems from the ground up, with active research support for extending the platform into clinical decision modeling.

TECHNICAL LEADERSHIP

Our team's work has been published in Nature journals, PNAS, JMIR, Chest, PLoS Computational Biology, The

Royal Society, and other leading venues.



Warren Pettine, MD, Co-Founder and CEO. Assistant Professor at the University of Utah where he leads the Medical Machine Intelligence (M²Int) Lab. Trained in machine learning research at Harvard, Stanford, NYU, and Yale. Prior health policy experience in the U.S. Congress and service on the University of Utah Institutional Review Board ground MTN's approach in policy and regulatory expertise.



Matthias Christenson, PhD, AI Architect. Investigator with the M²Int Lab. PhD and postdoctoral research at Columbia University in computational ML, with prior industry experience as a Deep Learning Research Engineer at DeepLife training foundational models on genomic and biometric data. Leads MTN's technical architecture design and data model development.



Brian Locke, MD, MSCI, Clinical AI Lead. Investigator with the M²Int Lab. Active ICU physician and Assistant Professor at Intermountain Healthcare, bringing firsthand understanding of clinical workflows across academic medical centers and integrated delivery networks. Provides the methodological rigor for the clinical and operational implications of MTN's technology.



Samuel Wecker, Lead Systems Engineer. Over twelve years building and scaling production software, including as a founding engineer at a startup that grew to a billion-dollar platform. Specializes in unifying disparate systems and data sources at scale. Leads Data Foundry's core platform development.

Contact:

Warren Pettine, MD, CEO of MTN

warren@themtn.ai

<https://www.themtn.ai>