TURN INTEGRATION LABOR INTO COMPOUNDING IP

Build the Accelerator Once. Compound It Across Every PE Bolt-On.

Executive Summary

The senior consultants on your B2B SaaS M&A team are exceptionally good at integrating disparate CRM, billing, and finance systems across PE-backed bolt-ons. The work is profitable, the methodology is real, and almost none of it carries to the next acquisition. Intellectual Property that should be accruing to your firm is being rebuilt by hand and walked out the door with each engagement.

The MTN Data Foundry augments your team's work with AI-assisted mapping that builds a semantic data model across sources. The Foundry operates on schema metadata: table and field names, data types, API documentation, internal data dictionaries, and your team's accumulated knowledge. Customer, contract, and revenue data values stay in your client's environment; the platform never sees them.

- **Your talent focuses on the high-value work.** Schema discovery and mapping that conventionally consumes 6 to 8 weeks compresses to 1 to 2 weeks during the pilot.
- **Maintain your data model, even as the software changes.** The Foundry automatically detects drift and suggests new mappings.
- **Every engagement deposits IP into your firm's library.** The Foundry captures three reusable assets on each deal: the annotation conventions your seniors apply, the firm's own documented mapping knowledge, and a growing library of canonical concepts (`Customer.AccountID`, `Contract.MRR`, `Opportunity.Stage`) that every future source maps against. Adding the tenth acquired CRM does not require touching the first nine. By engagement five, the marginal cost of integration is reviewing a queue, not building from scratch.

Intended Audience (Design Partner Briefing)

Target Audience: B2B SaaS M&A practice leaders evaluating a tool to deploy on the next PE bolt-on integration engagement.
Business Focus: A 90-day Design Partner pilot for the MTN Data Foundry. Single engagement, no exclusivity, mutually opt-in continuation.
Key Question: Can a productized integration platform cut our next engagement's CRM, billing, and finance data work and produce reusable templates for the bolt-on after that?

Comparison: Today's manual schema mapping across Salesforce, HubSpot, and NetSuite vs. an AI-assisted data fabric with B2B SaaS schema templates.

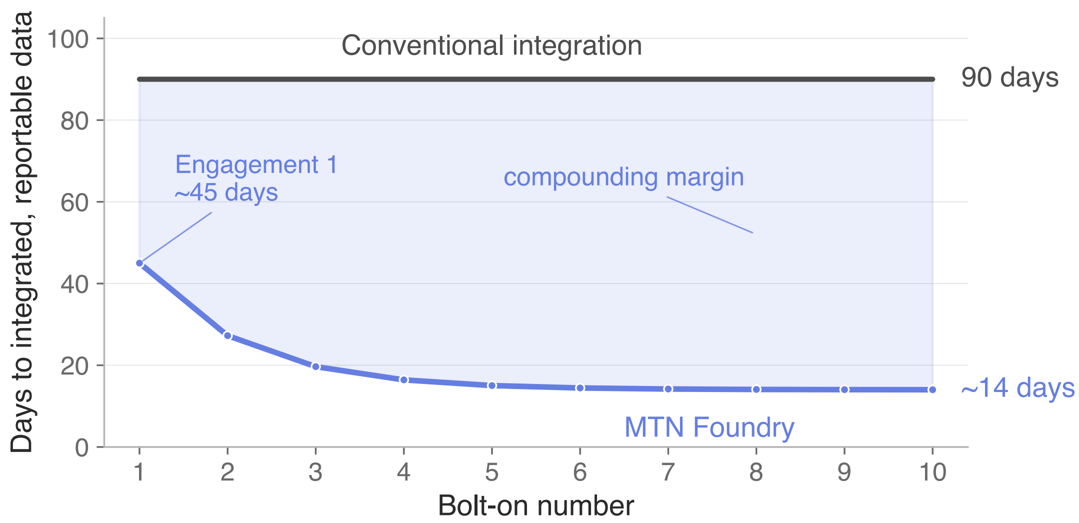

Figure 1. Days to integrated, reportable data per bolt-on under the conventional approach (flat at approximately three months) and under MTN Foundry. Two effects: engagement one is compressed by B2B SaaS templates and automated mapping (no compounding required); engagements two and onward compress further as cross-engagement reuse caches mapping decisions, declining toward a 14-day asymptote by the fifth bolt-on. Forward-looking model with conservative calibration.

Your Team's Talent is Under-Leveraged

Your senior consultants know the Salesforce metadata API, HubSpot's object model, and NetSuite's SuiteAnalytics segment dimensions at a level few in B2B SaaS PE can match. They read revenue-recognition edges and renewal-classification ambiguity that no generalist consultancy gets right. That expertise is what your firm's reputation rests on.

But B2B SaaS M&A integration is a slog. Three CRMs, two billing platforms, a stale Snowflake instance the platform CFO inherited, an operating partner asking for a unified ARR bridge across the platform by Friday. The pipelines you build break silently when Salesforce ships a winter release that renames a custom field, and none of the work carries to the next bolt-on.

The operating partner wants ARR by product line, NRR and GRR, CAC and payback, and pipeline coverage on a unified opportunity-stage taxonomy, all comparable across the platform and refreshed before the next LP letter. Nothing about that ask is new. What is new is that it has to land four times after four closes, then five, then six, on a clock your firm does not control.

The Foundry takes the mechanical mapping layer off your senior consultants' desks: schema discovery, field-level annotation, and the canonical concept map that lets one source plug into a shared semantic layer instead of into every other source pairwise. The judgment work stays human (architecture, exception handling, revenue-recognition and seat-counting calls), and value-level normalization stays where it belongs, in your firm's methodology and your client's data environment. The methodology your team applies is encoded once, audited, and carries into the next engagement.

A Solution that Converges

For the practice leader. Concurrency lift on every senior consultant, fixed-fee margin lift on every engagement that can shift off T&M, and a productized accelerator that survives consultant turnover.

For your senior integration architect. The mechanical mapping work disappears; architecture, exception handling, and revenue-recognition judgment stay human. Confidence scoring, audit logs, and the canonical schema model are visible and controllable.

For your engagement partner. A faster commitment your competitors cannot credibly match. Your team owns the client relationship and the deliverable; MTN engineering deploys in support, ensuring the platform performs against the engagement's specifics.

How the Foundry Works

The MTN Data Foundry is an AI-assisted data integration and normalization platform. Given a new source, it reasons over schema, software documentation, and your team's accumulated knowledge to discover the common data model and the mapping decisions that get there, with confidence scoring and human review on every key step.

The Foundry's integration technology operates on metadata: schema names, data types, API documentation, internal data dictionaries, and your team's accumulated knowledge. Customer, contract, and revenue data values stay in your client's environment; the platform never sees them. On top of that boundary, three AI-assisted capabilities replace the manual labor that consumes most of a conventional integration:

- **Annotations describing the function and intent of every table and field.** Your team spends hours hunting down the meaning of columns in someone else's Salesforce org or in a NetSuite instance whose segment, class, and department dimensions encode an undocumented chart of accounts. The Foundry reasons over database structure, software documentation, organizational documents, and your team's accumulated knowledge to annotate the database. Annotations are suggested by the system and flagged with confidence levels for human review.

- **Flexibly join data across sources.** Rather than building pairwise mappings between every source and every other source (the N×M explosion that consumes 65 to 75% of integration labor), the Foundry maps every source to a shared concept layer. Adding the tenth acquired CRM does not require touching the first nine; maintenance stays manageable as the portfolio grows.

- **Self-healing maintenance.** When Salesforce ships a Winter release that deprecates a Lead field, or NetSuite SuiteAnalytics renames a saved-search column, the Foundry detects the deviation, proposes the updated mapping, and queues it for review rather than letting the integration break silently weeks later.

The platform is bounded to the schema-and-concept layer by design. Value-level normalization (revenue-recognition rules across acquired billing systems, currency normalization, customer-account deduplication, product-SKU canonicalization) is a separately-deployable module that runs inside your client's environment when scope requires it. The metadata-driven approach extends to adjacent stacks (Microsoft Dynamics, Workday, Sage Intacct) where engagements require it.

What the Accelerator Is

An accelerator is an asset with pre-built tooling, methodology stacks, and AI-enabled platforms that layer into M&A integration projects. Every engagement you enter draws on the accumulated knowledge of engagements past. With the MTN Data Foundry, that asset has several constituent parts:

1. **Annotation conventions, your firm's house style, encoded.** Every engagement, senior consultants confirm or edit field-level annotations: usage notes, display names, definitions, and the reasoning behind each mapping. These are structured, versioned artifacts with confidence scores attached. The next engagement starts pre-loaded with the firm's accumulated conventions, and the suggestion engine is calibrated against the firm's own gold-standard reviews.
2. **Internal documented knowledge, made operational.** The Foundry reasons over your firm's own integration playbooks, methodology decks, prior engagement deliverables, and internal data dictionaries, using them as context for mapping reasoning so first-pass annotations reflect your documented approach rather than a generic prior. Document libraries that previously lived dormant in SharePoint become active mapping intelligence on every engagement.
3. **Canonical concepts derived from past engagements.** A growing library of stable business objects (`Customer.AccountID`, `Contract.MRR`, `Contract.RenewalDate`, `Subscription.SeatCount`, `Opportunity.Stage`, `Product.SKU`) that every future source maps against. By engagement five, the library has crystallized around your firm's B2B SaaS roll-up patterns (CRM consolidation, billing-platform

migration, ARR bridge construction, customer-360 unification), and the marginal cost of the next bolt-on falls toward the cost of reviewing a queue.

4. **Vendor-schema drift signatures and reviewer feedback loops.** When Salesforce, HubSpot, NetSuite, or any of the major billing platforms ships a release that renames a field, the Foundry's correction is captured and applies firm-wide. Every approve, reject, or edit decision a senior consultant makes becomes labeled training data calibrating the system for your firm's standards.

The MTN Data Foundry accelerates work during the pilot period, and value compounds with every subsequent engagement. Pilot benchmarks suggest a senior consultant who today delivers 4 to 6 PE bolt-on integrations per year can deliver materially more on a productized platform, at higher margin per deal. We will commit to specific targets jointly with our design partner based on engagement scope.

What Engagement One Looks Like

The conventional first integration on a typical B2B SaaS PE bolt-on takes roughly three months. The Foundry compresses that work substantially on engagement one, before any cross-engagement compounding has kicked in. Engagement one delivers a canonical concept map across the in-scope CRM, billing, and finance instances, and a queryable semantic layer your team can build the operating partner's ARR bridge, NRR/GRR, and pipeline-coverage dashboards on top of. MTN's engineering talent deploys in support of your team, ensuring the Data Foundry delivers.

What Success Looks Like

A successful pilot produces:

- Detailed annotations for every table and field, facilitating both general business intelligence and the use of advanced AI tooling.
- A clear mapping of data sources to either a target schema, or a canonical set of concepts.
- A unified ARR, retention, and pipeline dashboard delivered to the operating partner, built on the canonical concept map the Foundry produces, with your firm's normalization and revenue-recognition rules applied in your delivery environment.
- An audit log of every transform applied.

After the pilot, the relationship is mutually opt-in. Whether you continue or not, you keep your annotation conventions, methodology playbooks, vertical templates, and engagement deliverables, exported as portable JSON manifests. MTN keeps the platform and the learnings about how it performs across deployments. License terms, support model, and account scoping for an ongoing relationship are documented in the design partner term sheet, available on request.

What We Are Asking For

We are seeking our first B2B SaaS M&A integration design partner. The terms below reflect that: discounted, scope-bounded, and structured to build a case study together.

Three concrete asks over the next 90 days.

1. A 30-minute scoping conversation with the M&A practice leader for B2B SaaS and the engagement lead for an upcoming PE bolt-on, to confirm whether the Foundry's current CRM, billing, and finance coverage matches the engagement.
2. Identification of a candidate engagement in your firm's pipeline (CRM consolidation, billing-platform migration, ARR-bridge construction, vertical SaaS aggregator integration, horizontal SaaS tuck-in, SaaS carve-out) that could host the 90-day pilot.
3. A 90-day Design Partner agreement covering platform license discount, engineering support, and other key details. Draft term sheet ready for review.

No exclusive in either direction.

For the platform's technical architecture and compliance posture, request the *MTN Data Foundry Architecture and Compliance Brief*.

About MTN

MTN is a team with deep expertise in data infrastructure, machine learning, and semantic data systems. The founding team has built schema-mapping and data-integration systems from the ground up at production scale.

Technical Leadership

Warren Pettine, MD, Co-Founder and CEO. Assistant Professor at the University of Utah, where he leads the Medical Machine Intelligence (M^2Int) Lab. Trained in machine learning research at Harvard, Stanford, NYU, and Yale.

Matthias Christenson, PhD, AI Architect. Leads MTN's technical architecture design and data model development. PhD and postdoctoral research at Columbia University in computational machine learning. Prior industry experience as a Deep Learning Research Engineer at DeepLife, training foundational models on genomic and biometric data.

Samuel Wecker, Lead Systems Engineer. Over twelve years building and scaling production software, including as a founding engineer at a startup that grew to a billion-dollar platform. Specializes in unifying disparate systems and data sources at scale. Leads Data Foundry's core platform development.

Contact:
Warren Pettine, MD, CEO of MTN
warren@themtn.ai
https://www.themtn.ai